

Alex Johnson

alex.johnson@example.com | (555) 234-5678 | San Francisco, CA | linkedin.com/in/alexjohnson-ml

SUMMARY

Lead Machine Learning Engineer with 6+ years designing large-scale NLP, PyTorch transformers, multimodal perception systems, and distributed model inference services in cloud environments.

EDUCATION

Stanford University - Bachelor of Science in Computer Science (2016 - 2020)

Focus: Artificial Intelligence, Deep Learning, and Mathematical Statistics

SKILLS

Technical Skills: Python, PyTorch, TensorFlow, Scikit-learn, Transformers, CUDA, HuggingFace

Systems: Docker, Kubernetes, AWS, FastAPI, Ray, Vector Databases, PostgreSQL, Git

Specializations: Natural Language Processing, Model Optimization, RLHF, Embeddings, MLOps

EXPERIENCE

Staff Machine Learning Engineer - Scale AI (2022 - Present)

- Led fine-tuning and deployment of vision and language foundation models for enterprise workloads.
- Architected high-throughput inference microservices reducing p99 latency by 45%.
- Built automated continuous evaluation harnesses measuring vector alignment and model safety.

Senior Machine Learning Engineer - Databricks (2020 - 2022)

- Scaled distributed PyTorch training jobs over multi-node GPU clusters with Ray and Spark.
- Implemented feature store transformations and automated ETL pipelines for petabyte datasets.

PROJECTS

LLM Alignment & Evaluation Engine: Open-source preference optimization harness for frontier LLMs.

Real-Time Multimodal Search: Semantic vector similarity search engine indexing 50M+ items.